# Income Segregation and Intergenerational Mobility Across Colleges in the United States

Jacob Sharadin

# Paper Background

Main idea: Investigating how higher education shapes mobility in the US.

How could changes in the distributions of students from different backgrounds in college affect segregation by parental income and intergenerational mobility?

# Paper Background

Paper Outline - Three Main Sections:

1. Constructing statistics on parent income at different colleges
2. Earning outcomes of students at each college
3. Changes in income segregation if students were allocated to colleges evenly

# Paper Findings

Four Key Findings:

1. Low and middle income students attend selective schools at much lower rates than students from higher income
2. Very low middle class representation at most selective schools (Ivy-Plus)
3. Would need to raise attendance rates for low income students from ~7% to ~26%
    a. Low-income students would need to attend all schools at rates similar to those with 160 point higher SAT scores
4. By equalizing attendance rates for students with the same test scores, outcome gap would be decreased by 15%

Main takeaway: By changing how students are allocated to colleges, segregation could be decreased and intergenerational mobility increased

# The Dataset

- College attendance: federal tax records, Department of Education records 1999-2013
- Incomes: federal income tax 1996-2014, information returns (like W-2)
- Parent Income: total pre-tax income at the household level
  - Averaged over the five years when child aged 15-19
  - Parents then assigned income percentiles through ranking with other parents w/ children in same birth cohort
- Child Income: pre-tax individual earnings 2014
  - Ranked relative to other children in same birth cohort
- ~2200 observations

# Predicting k_rank

```
Number of categories in k_rank: 2199
Number of categories in count: 1773
Number of categories in female: 2160
Number of categories in par_mean: 2199
Number of categories in par_median: 910
Number of categories in par_rank: 2199
Number of categories in type: 3
Number of categories in tier_name: 12
Number of categories in iclevel: 3
Number of categories in region: 4
```

- Numerical variables
  - k_rank (Dependent Variable)
  - count
  - female
  - par_mean
  - par_median
  - par_rank
- Categorical variables
  - type
  - tier_name
  - iclevel
  - region

# Handling Categorical Variables

| | |
|---|---|
| type | Type :<br>1 = public<br>2 = private non-profit<br>3 = for-profit |
| tier | Selectivity and type combination (see Table 6 for more detailed descriptions of these groups):<br>1 = Ivy Plus<br>2 = Other elite schools (public and private)<br>3 = Highly selective public<br>4 = Highly selective private<br>5 = Selective public<br>6 = Selective private<br>7 = Nonselective 4-year public<br>8 = Nonselective 4-year private not-for-profit<br>9 = Two-year (public and private not-for-profit)<br>10 = Four-year for-profit<br>11 = Two-year for-profit<br>12 = Less than two year schools of any type<br>13 = Attending college with insufficient data<br>14 = Not in college between the ages of 19-22 |
| tier_name | Name of college tier |

One-hot encoding

| type_for-profit | type_private non-profit | type_public | tier_name_Four-year for-profit | tier_name_Highly selective private | tier_name_Highly selective public | ... |
|---|---|---|---|---|---|---|
| 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... |
| 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... |
| 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... |
| 1.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | ... |
| 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... |

# Data Cleaning

```
k_rank          0
count           0
female         19
par_mean        0
par_median      0
par_rank        0
type            0
tier_name       0
iclevel         0
region          0
dtype: int64
```

- Only 19 missing values

# Summary Statistics and Boxplots

| | k_rank | count | female | par_mean | par_median | par_rank |
|---|---|---|---|---|---|---|
| **count** | 2202.000000 | 2202.000000 | 2183.000000 | 2202.000000 | 2202.000000 | 2202.000000 |
| **mean** | 0.567720 | 1714.291023 | 0.555279 | 107432.511713 | 77695.458674 | 0.572805 |
| **std** | 0.086629 | 23243.749136 | 0.139493 | 67386.449844 | 28463.280143 | 0.117411 |
| **min** | 0.340474 | 50.000000 | 0.003306 | 33202.243485 | 21200.000000 | 0.252361 |
| **25%** | 0.506592 | 232.000000 | 0.504596 | 69841.513082 | 59100.000000 | 0.489515 |
| **50%** | 0.554700 | 467.583333 | 0.550342 | 88621.716206 | 74300.000000 | 0.574253 |
| **75%** | 0.626928 | 1038.333333 | 0.599742 | 118488.889985 | 91700.000000 | 0.655498 |
| **max** | 0.906024 | 955065.333333 | 1.000000 | 551968.154148 | 226700.000000 | 0.887999 |

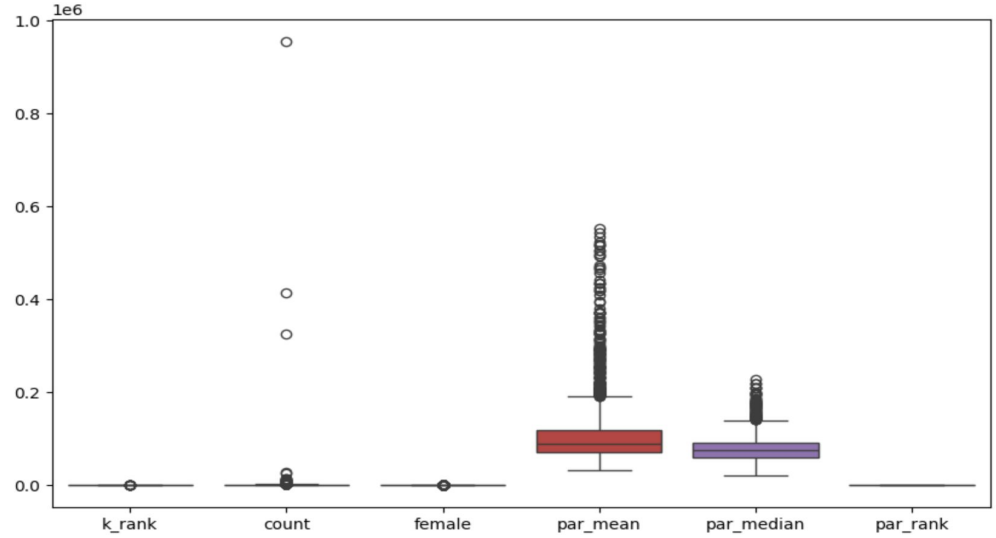

# Identifying Highly Correlated Variables

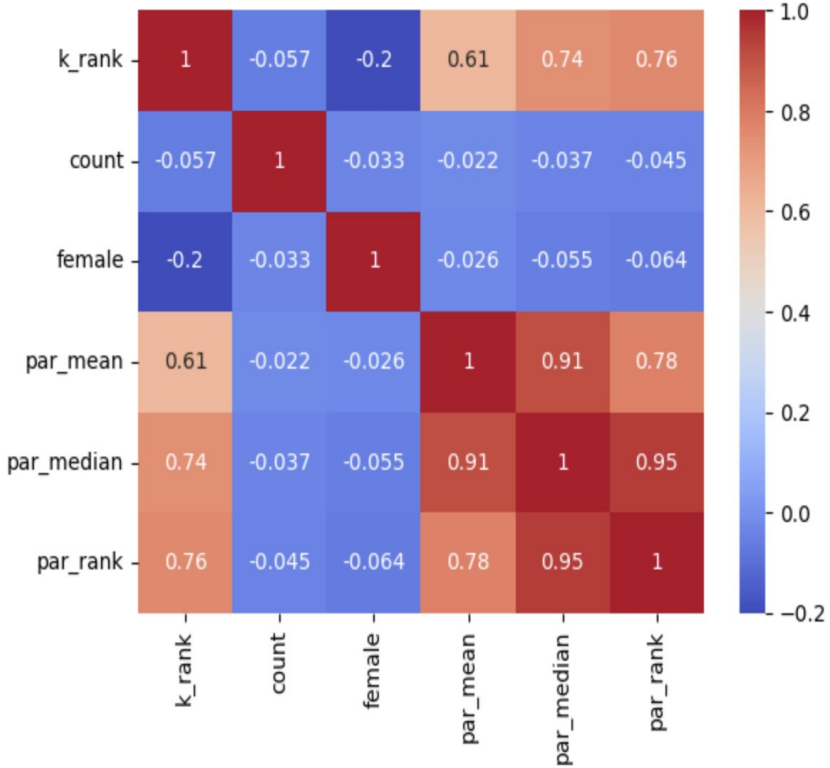


- Multicollinearity between variables

# Pairplot Visualization

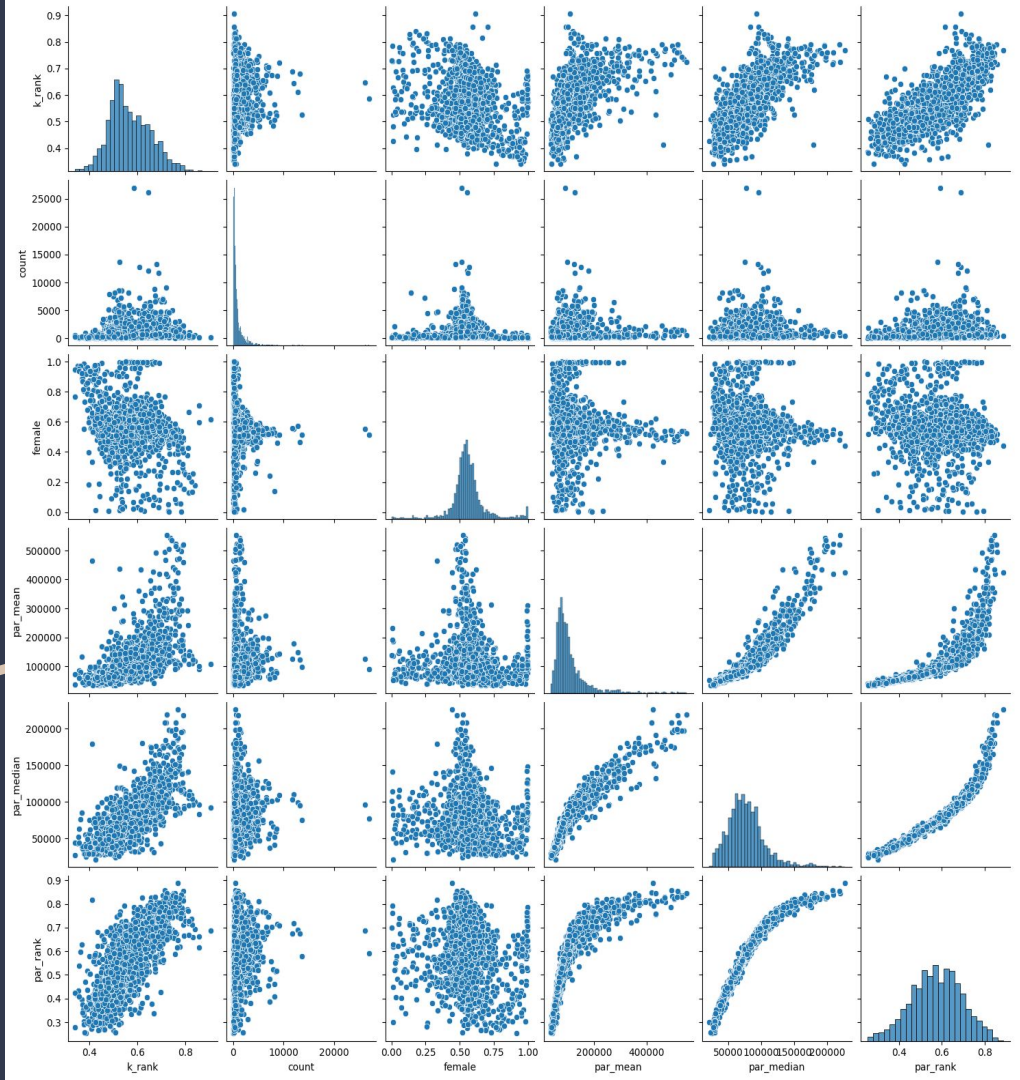

# Regression Analysis

OLS Regression Results

| | | | |
|---|---|---|---|
| Dep. Variable: | k_rank | R-squared: | 0.727 |
| Model: | OLS | Adj. R-squared: | 0.725 |
| Method: | Least Squares | F-statistic: | 338.5 |
| Date: | Wed, 26 Jun 2024 | Prob (F-statistic): | 0.00 |
| Time: | 22:36:44 | Log-Likelihood: | 3657.5 |
| No. Observations: | 2180 | AIC: | -7279. |
| Df Residuals: | 2162 | BIC: | -7177. |
| Df Model: | 17 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P>\|t\| | [0.025 | 0.975] |
|---|---|---|---|---|---|---|
| Intercept | 0.3555 | 0.008 | 43.226 | 0.000 | 0.339 | 0.372 |
| par_rank | 0.3147 | 0.012 | 26.522 | 0.000 | 0.291 | 0.338 |
| type_private_non_profit | 0.0434 | 0.021 | 2.045 | 0.041 | 0.002 | 0.085 |
| type_public | 0.0456 | 0.020 | 2.248 | 0.025 | 0.006 | 0.085 |
| tier_name_Highly_selective_private | 0.0450 | 0.023 | 1.975 | 0.048 | 0.000 | 0.090 |
| tier_name_Highly_selective_public | 0.0900 | 0.023 | 3.913 | 0.000 | 0.045 | 0.135 |
| tier_name_Ivy_Plus | 0.1075 | 0.026 | 4.158 | 0.000 | 0.057 | 0.158 |
| tier_name_Less_than_two_year_schools_of_any_type | -0.0381 | 0.005 | -8.463 | 0.000 | -0.047 | -0.029 |
| tier_name_Nonselective_four_year_private_not_for_profit | -0.0454 | 0.023 | -2.016 | 0.044 | -0.090 | -0.001 |
| tier_name_Nonselective_four_year_public | -0.0304 | 0.022 | -1.404 | 0.160 | -0.073 | 0.012 |
| tier_name_Other_elite_schools_public_and_private | 0.0753 | 0.023 | 3.317 | 0.001 | 0.031 | 0.120 |
| tier_name_Selective_private | 0.0204 | 0.022 | 0.926 | 0.355 | -0.023 | 0.064 |
| tier_name_Selective_public | 0.0245 | 0.021 | 1.160 | 0.246 | -0.017 | 0.066 |
| tier_name_Two_year_public_and_private_not_for_profit | -0.0274 | 0.014 | -1.987 | 0.047 | -0.055 | -0.000 |
| tier_name_Two_year_for_profit | 0.0090 | 0.008 | 1.129 | 0.259 | -0.007 | 0.024 |
| iclevel_Less_than_Two_year | -0.0381 | 0.005 | -8.463 | 0.000 | -0.047 | -0.029 |
| iclevel_Two_year | -0.0185 | 0.008 | -2.365 | 0.018 | -0.034 | -0.003 |
| region_Northeast | 0.0136 | 0.003 | 4.707 | 0.000 | 0.008 | 0.019 |
| region_South | -0.0142 | 0.003 | -5.357 | 0.000 | -0.019 | -0.009 |
| region_West | -0.0123 | 0.003 | -3.943 | 0.000 | -0.018 | -0.006 |

| | | | |
|---|---|---|---|
| Omnibus: | 392.534 | Durbin-Watson: | 1.814 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 2159.379 |
| Skew: | 0.734 | Prob(JB): | 0.00 |
| Kurtosis: | 7.649 | Cond. No. | 3.08e+16 |

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 5.48e-30. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.

- This result suggests that this model is not valid

# Prediction with the use of ML algorithms

Model #1: Include all features to get initial results

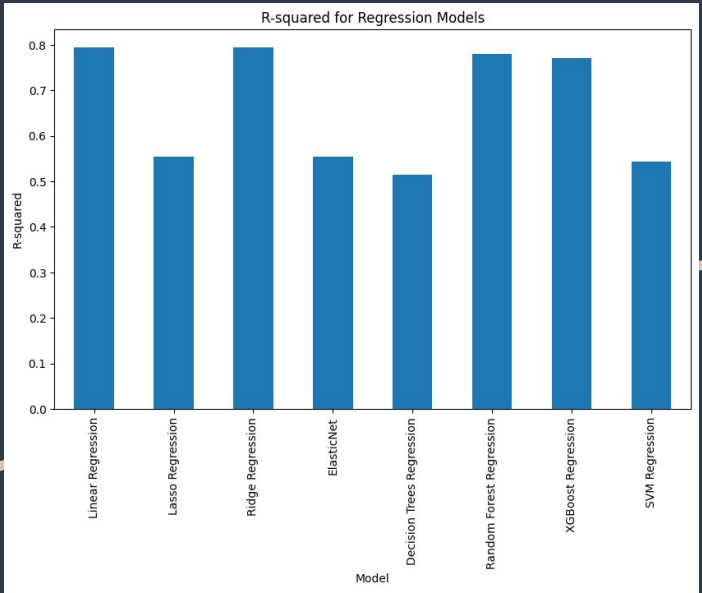


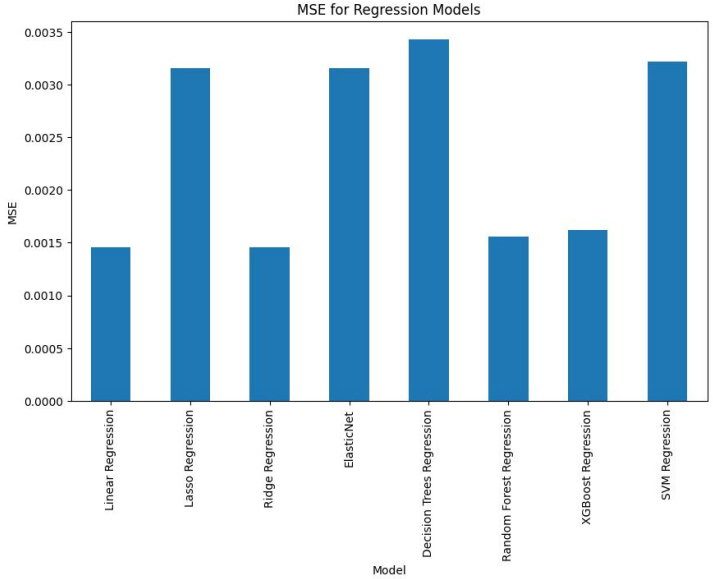

# Model #1 Feature Importances

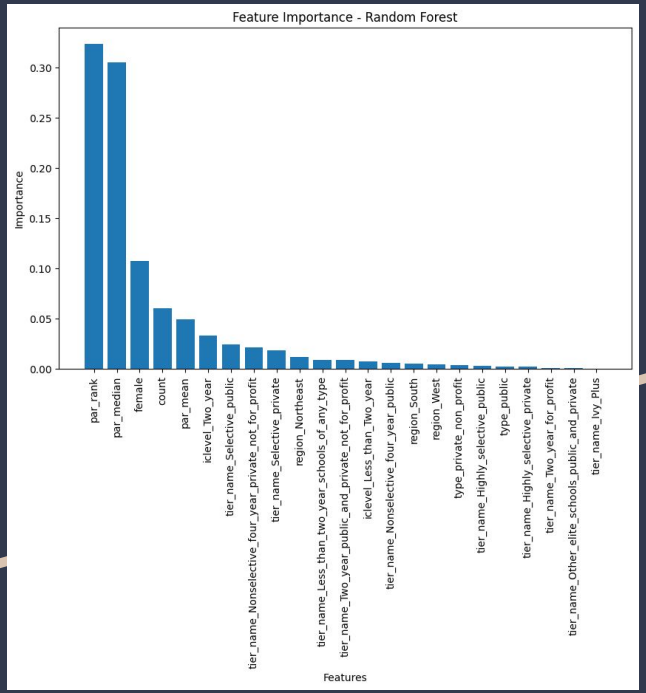


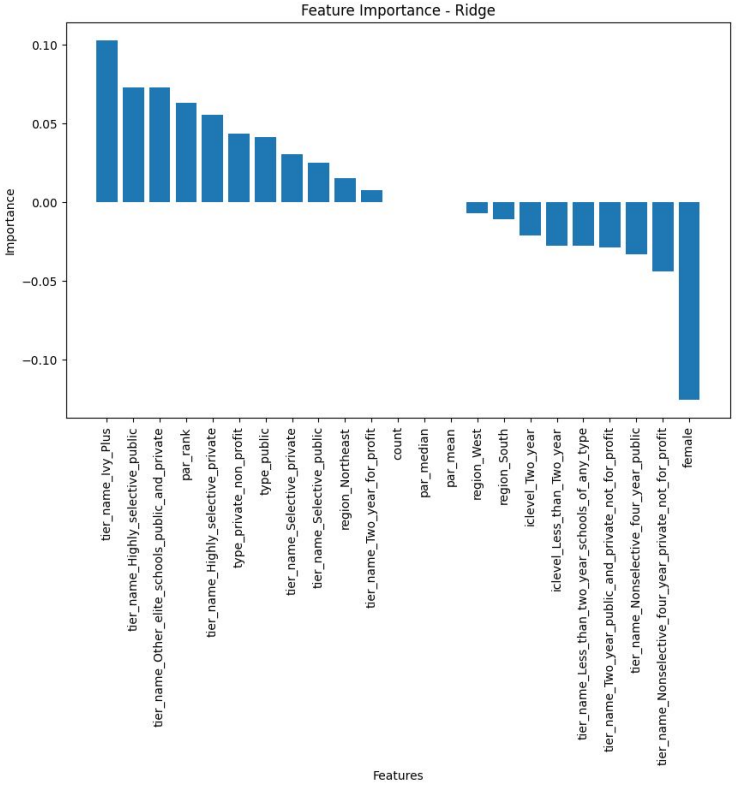

# Prediction Model #2

Model #2: Drop par_mean and par_median

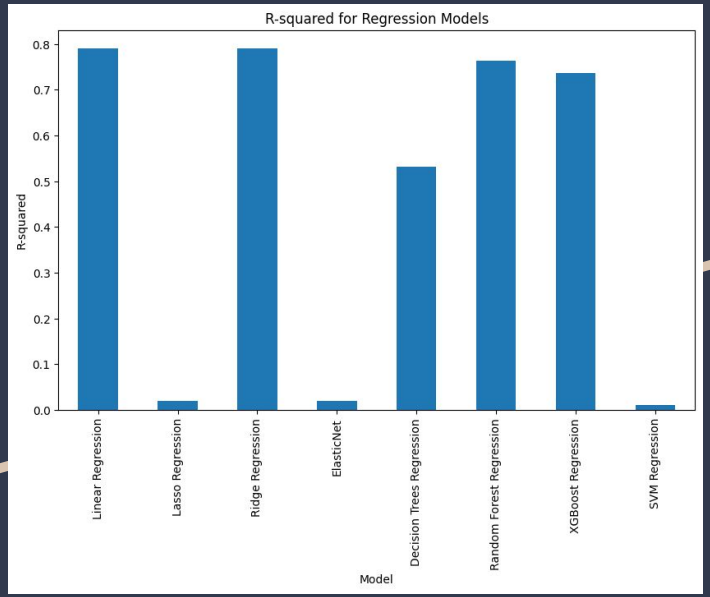


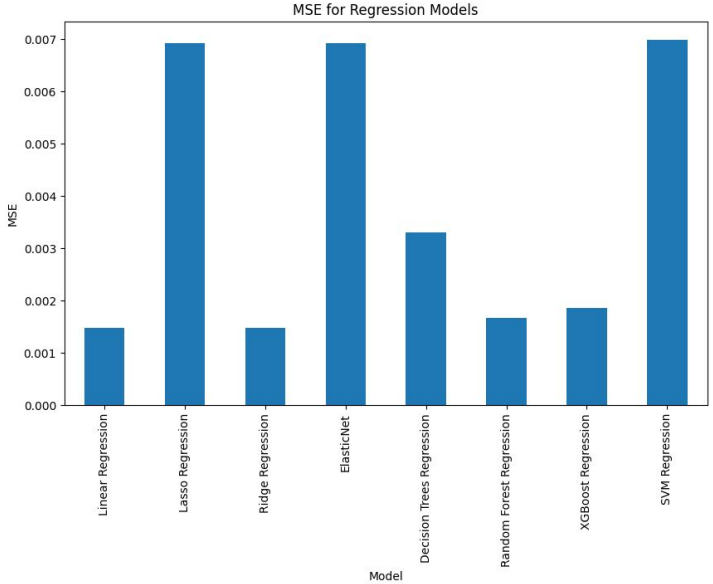

# Model #2 Feature Importances

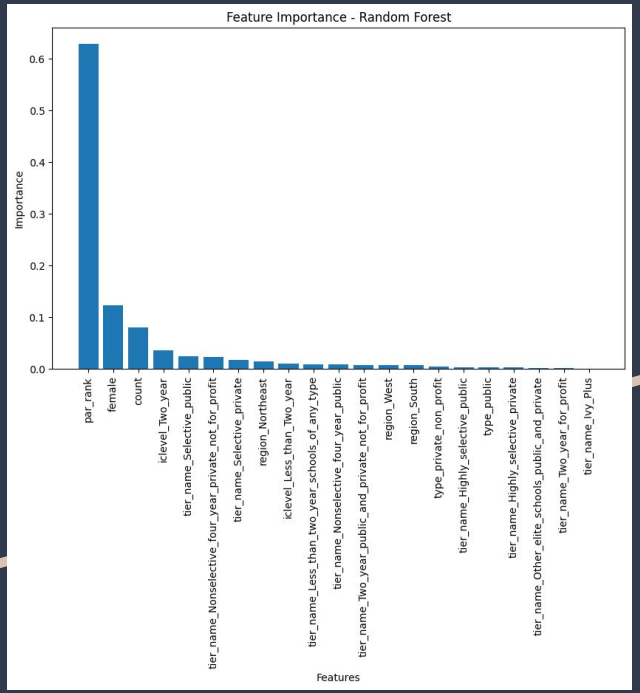


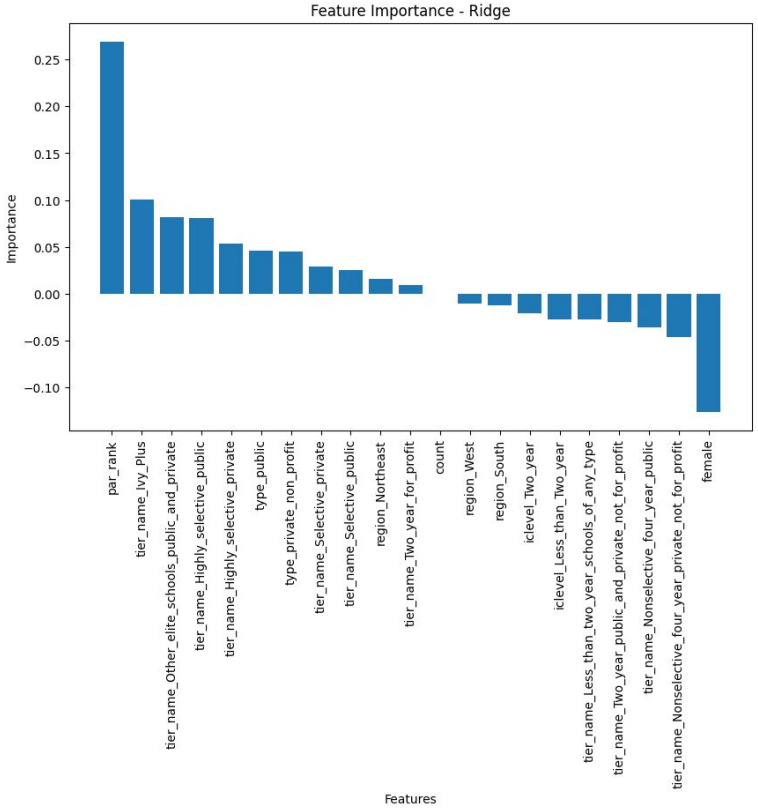

# Ridge Regression Feature Importances

| Feature | Importance |
|---|---|
| par_rank | 0.269273 |
| female | 0.125790 |
| tier_name_Ivy_Plus | 0.100166 |
| tier_name_Other_elite_schools_public_and_private | 0.081286 |
| tier_name_Highly_selective_public | 0.081103 |
| tier_name_Highly_selective_private | 0.053809 |
| tier_name_Nonselective_four_year_private_not_f... | 0.046604 |
| type_public | 0.045542 |
| type_private_non_profit | 0.044869 |
| tier_name_Nonselective_four_year_public | 0.036212 |
| tier_name_Two_year_public_and_private_not_for_... | 0.030178 |
| tier_name_Selective_private | 0.029270 |
| tier_name_Less_than_two_year_schools_of_any_type | 0.027441 |
| iclevel_Less_than_Two_year | 0.027441 |
| tier_name_Selective_public | 0.024941 |
| iclevel_Two_year | 0.021048 |
| region_Northeast | 0.015683 |
| region_South | 0.012488 |
| region_West | 0.010901 |
| tier_name_Two_year_for_profit | 0.009130 |
| count | 0.000002 |

Absolute value taken for each feature importance value

# Prediction Model #3

Model #3: Keeping the 10 most important features

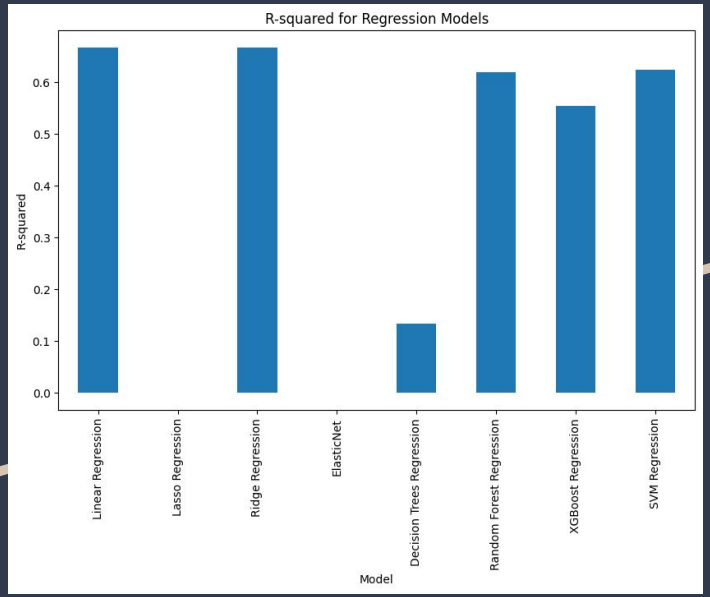


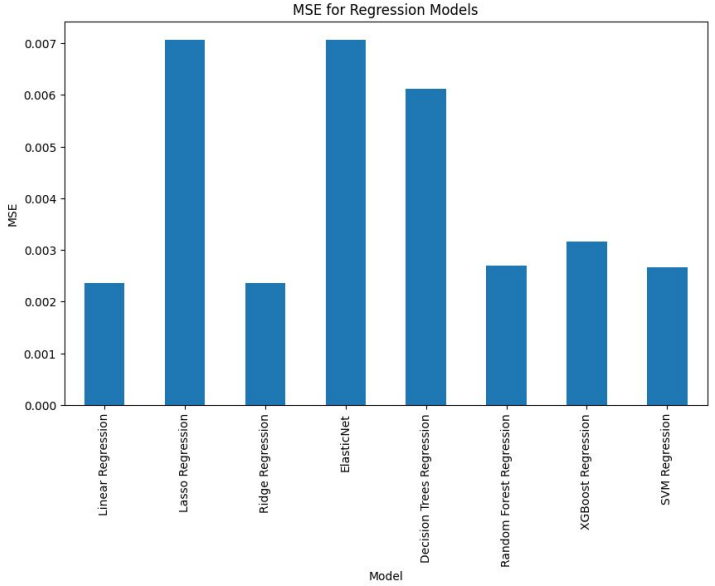

# Interpretations

Results for Train-Test Split:

| Model | R-squared | Mean Squared Error | Root Mean Squared Error |
|---|---|---|---|
| Linear Regression (Train-Test Split) | 0.6956 | 0.0021 | 0.0453 |
| Lasso (Train-Test Split) | -0.0002 | 0.0068 | 0.0822 |
| ElasticNet (Train-Test Split) | -0.0002 | 0.0068 | 0.0822 |
| Decision Tree (Train-Test Split) | 0.3877 | 0.0041 | 0.0643 |
| XGBoost (Train-Test Split) | 0.6052 | 0.0027 | 0.0516 |
| Ridge (Train-Test Split) | 0.6947 | 0.0021 | 0.0454 |
| SVM (Train-Test Split) | 0.6325 | 0.0025 | 0.0498 |
| Random Forest (Train-Test Split) | 0.6569 | 0.0023 | 0.0481 |

Results for Whole Data:

| Model | R-squared | Mean Squared Error | Root Mean Squared Error |
|---|---|---|---|
| Linear Regression (Whole Data) | 0.6669 | 0.0025 | 0.0499 |
| Lasso (Whole Data) | 0.0000 | 0.0075 | 0.0865 |
| ElasticNet (Whole Data) | 0.0000 | 0.0075 | 0.0865 |
| Decision Tree (Whole Data) | 1.0000 | 0.0000 | 0.0000 |
| XGBoost (Whole Data) | 0.9363 | 0.0005 | 0.0218 |
| Ridge (Whole Data) | 0.6664 | 0.0025 | 0.0500 |
| SVM (Whole Data) | 0.6287 | 0.0028 | 0.0527 |
| Random Forest (Whole Data) | 0.9501 | 0.0004 | 0.0193 |

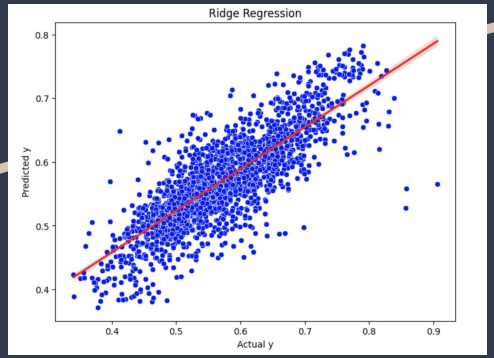


- While the Random Forest model has an r squared value of 0.9501 and decision tree of 1.0000. Ridge model performs best on new data.
- Ridge Regression is the model we selected to characterize our dataset.
- Feature importance (top 10)
- Results from OLS regression with top 10 features
- Scatterplot of ridge regression model with line of best fit

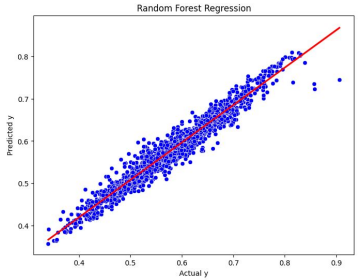

# Final Results

- Regression results
  - Using top 10 features from Ridge

OLS Regression Results

| | | | |
|---|---|---|---|
| Dep. Variable: | k_rank | R-squared: | 0.771 |
| Model: | OLS | Adj. R-squared: | 0.769 |
| Method: | Least Squares | F-statistic: | 345.6 |
| Date: | Thu, 27 Jun 2024 | Prob (F-statistic): | 0.00 |
| Time: | 08:14:54 | Log-Likelihood: | 3848.4 |
| No. Observations: | 2180 | AIC: | -7653. |
| Df Residuals: | 2158 | BIC: | -7528. |
| Df Model: | 21 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P>\|t\| | [0.025 | 0.975] |
|---|---|---|---|---|---|---|
| Intercept | 0.4673 | 0.011 | 41.208 | 0.000 | 0.445 | 0.490 |
| count | 1.604e-06 | 6.78e-07 | 2.366 | 0.018 | 2.75e-07 | 2.93e-06 |
| female | -0.1292 | 0.007 | -19.024 | 0.000 | -0.143 | -0.116 |
| par_mean | -3.139e-07 | 4.64e-08 | -6.760 | 0.000 | -4.05e-07 | -2.23e-07 |
| par_median | 1.271e-06 | 2.19e-07 | 5.805 | 0.000 | 8.42e-07 | 1.7e-06 |
| par_rank | 0.1187 | 0.037 | 3.187 | 0.001 | 0.046 | 0.192 |
| type_private_non_profit | 0.0079 | 0.020 | 0.402 | 0.688 | -0.031 | 0.046 |
| type_public | 0.0115 | 0.019 | 0.613 | 0.540 | -0.025 | 0.048 |
| tier_name_Highly_selective_private | 0.0984 | 0.021 | 4.666 | 0.000 | 0.057 | 0.140 |
| tier_name_Highly_selective_public | 0.1115 | 0.021 | 5.242 | 0.000 | 0.070 | 0.153 |
| tier_name_Ivy_Plus | 0.1609 | 0.025 | 6.502 | 0.000 | 0.112 | 0.209 |
| tier_name_Less_than_two_year_schools_of_any_type | -0.0237 | 0.004 | -5.653 | 0.000 | -0.032 | -0.016 |
| tier_name_Nonselective_four_year_private_not_for_profit | -0.0056 | 0.021 | -0.270 | 0.787 | -0.046 | 0.035 |
| tier_name_Nonselective_four_year_public | -0.0011 | 0.020 | -0.055 | 0.956 | -0.040 | 0.038 |
| tier_name_Other_elite_schools_public_and_private | 0.1200 | 0.021 | 5.651 | 0.000 | 0.078 | 0.162 |
| tier_name_Selective_private | 0.0695 | 0.020 | 3.403 | 0.001 | 0.029 | 0.110 |
| tier_name_Selective_public | 0.0581 | 0.019 | 2.980 | 0.003 | 0.020 | 0.096 |
| tier_name_Two_year_public_and_private_not_for_profit | -0.0081 | 0.013 | -0.634 | 0.526 | -0.033 | 0.017 |
| tier_name_Two_year_for_profit | -0.0010 | 0.007 | -0.139 | 0.889 | -0.015 | 0.013 |
| iclevel_Less_than_Two_year | -0.0237 | 0.004 | -5.653 | 0.000 | -0.032 | -0.016 |
| iclevel_Two_year | -0.0091 | 0.007 | -1.263 | 0.207 | -0.023 | 0.005 |
| region_Northeast | 0.0158 | 0.003 | 5.926 | 0.000 | 0.011 | 0.021 |
| region_South | -0.0105 | 0.002 | -4.241 | 0.000 | -0.015 | -0.006 |
| region_West | -0.0070 | 0.003 | -2.381 | 0.017 | -0.013 | -0.001 |

| | | | |
|---|---|---|---|
| Omnibus: | 375.416 | Durbin-Watson: | 1.770 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 3193.556 |
| Skew: | 0.556 | Prob(JB): | 0.00 |
| Kurtosis: | 8.824 | Cond. No. | 3.79e+21 |

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 3.44e-30. This might indicate that there are

OLS Regression Results

| | | | |
|---|---|---|---|
| Dep. Variable: | k_rank | R-squared: | 0.667 |
| Model: | OLS | Adj. R-squared: | 0.666 |
| Method: | Least Squares | F-statistic: | 434.8 |
| Date: | Thu, 27 Jun 2024 | Prob (F-statistic): | 0.00 |
| Time: | 08:15:02 | Log-Likelihood: | 3441.9 |
| No. Observations: | 2180 | AIC: | -6862. |
| Df Residuals: | 2169 | BIC: | -6799. |
| Df Model: | 10 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P>\|t\| | [0.025 | 0.975] |
|---|---|---|---|---|---|---|
| Intercept | 0.3576 | 0.009 | 41.749 | 0.000 | 0.341 | 0.374 |
| female | -0.1148 | 0.008 | -14.503 | 0.000 | -0.130 | -0.099 |
| par_rank | 0.4077 | 0.012 | 34.424 | 0.000 | 0.384 | 0.431 |
| type_private_non_profit | 0.0679 | 0.005 | 14.745 | 0.000 | 0.059 | 0.077 |
| type_public | 0.0277 | 0.004 | 6.809 | 0.000 | 0.020 | 0.036 |
| tier_name_Highly_selective_private | 0.0130 | 0.006 | 2.013 | 0.044 | 0.000 | 0.026 |
| tier_name_Highly_selective_public | 0.0937 | 0.010 | 9.252 | 0.000 | 0.074 | 0.114 |
| tier_name_Ivy_Plus | 0.0695 | 0.015 | 4.711 | 0.000 | 0.041 | 0.098 |
| tier_name_Nonselective_four_year_private_not_for_profit | -0.0686 | 0.006 | -11.243 | 0.000 | -0.081 | -0.057 |
| tier_name_Nonselective_four_year_public | -0.0135 | 0.006 | -2.214 | 0.027 | -0.025 | -0.002 |
| tier_name_Other_elite_schools_public_and_private | 0.0418 | 0.007 | 6.122 | 0.000 | 0.028 | 0.055 |

| | | | |
|---|---|---|---|
| Omnibus: | 256.542 | Durbin-Watson: | 1.698 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 901.434 |
| Skew: | 0.561 | Prob(JB): | 1.80e-196 |
| Kurtosis: | 5.943 | Cond. No. | 20.9 |

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

# Some Drawbacks and Improvements

- More extensive hyperparameter tuning could've been done
  - Used alpha value = 0.1
- Selection of top 10 features
  - Ex. In model 3, choose top 10 features
  - Somewhat arbitrary, could've chosen 9, 11, etc.

# Conclusion

# Works Cited

1. Raj Chetty, John N Friedman, Emmanuel Saez, Nicholas Turner, Danny Yagan, Income Segregation and Intergenerational Mobility Across Colleges in the United States, *The Quarterly Journal of Economics*, Volume 135, Issue 3, August 2020, Pages 1567–1633, https://doi.org/10.1093/qje/qjaa005